

The Efficient Frontier: Markowitz portfolio optimization in Python

Posted by: Thomas Wiecki

[« Previous Post](#)

[Next Post »](#)

Authors: [Dr. Thomas Starke](#), David Edwards, Dr. Thomas Wiecki

Introduction

In this blog post you will learn about the basic idea behind Markowitz portfolio optimization as well as how to do it in Python. We will then show how you can create a simple backtest that rebalances its portfolio in a Markowitz-optimal way. We hope you enjoy it and get a little more enlightened in the process.

We will start by using random data and only later use actual stock data. This will hopefully help you to get a sense of how to use modelling and simulation to improve your understanding of the theoretical concepts. Don't forget that the skill of an algo-trader is to put mathematical models into code and this example is great practice.

Let's start with importing a few modules, which we need later and produce a series of normally distributed returns. `cvxopt` is a convex solver which you can easily download with `sudo pip install cvxopt`.

Simulations

```
%matplotlib inline
import numpy as np
import matplotlib.pyplot as plt
import cvxopt as opt
from cvxopt import blas, solvers
import pandas as pd

np.random.seed(123)

# Turn off progress printing
solvers.options['show_progress'] = False
```

Assume that we have 4 assets, each with a return series of length 1000. We can use `numpy.random.randn` to sample returns from a normal distribution.

```
## NUMBER OF ASSETS
n_assets = 4

## NUMBER OF OBSERVATIONS
n_obs = 1000

return_vec = np.random.randn(n_assets, n_obs)
```

```
plt.plot(return_vec.T, alpha=.4);
plt.xlabel('time')
plt.ylabel('returns')
```

```
<matplotlib.text.Text at 0x7fa919b8c690>
```

Follow Us!



Quantopian

[Learn More About Quantopian](#)

[Watch A Video About Quantopian](#)

Get email updates about Quantopian:

<input type="text" value="Enter email"/>	<input type="button" value="Subscribe"/>
--	--

Recent Posts

[The Q: Announcing Our New Logo](#)

[Quantopian's First Discretionary Capital Allocations](#)

[Bayesian Optimization of a Technical Trading Algorithm with Zipline+SigOpt](#)

 [Get Blog Updates by Email](#)

Recent Comments

Roger Quiring on [The Q: Announcing Our New Logo](#)

Chris on [Quantopian's First Discretionary Capital Allocations](#)

Scott Clark on [Bayesian Optimization of a Technical Trading Algorithm with Zipline+SigOpt](#)

Categories

[Data Transform](#)

[Interview](#)

[Optimization](#)

[QuantCon](#)

[Quantopian](#)

Archives

[March 2016](#)

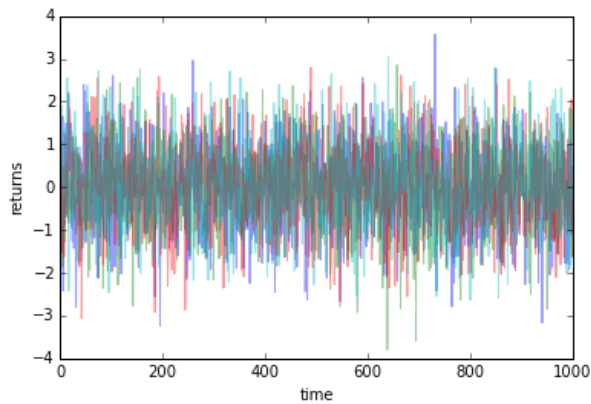
[February 2016](#)

[January 2016](#)

[November 2015](#)

[October 2015](#)

[September 2015](#)



These return series can be used to create a wide range of portfolios, which all have different returns and risks (standard deviation). We can produce a wide range of random weight vectors and plot those portfolios. As we want all our capital to be invested, this vector will have to sum to one.

```
def rand_weights(n):
    ''' Produces n random weights that sum to 1 '''
    k = np.random.rand(n)
    return k / sum(k)

print rand_weights(n_assets)
print rand_weights(n_assets)
```

```
[ 0.54066805  0.2360283  0.11660484  0.1066988 ]
[ 0.27638339  0.03006307  0.47850085  0.21505269 ]
```

Next, let's evaluate how many of these random portfolios would perform. Towards this goal we are calculating the mean returns as well as the volatility (here we are using standard deviation). You can also see that there is a filter that only allows to plot portfolios with a standard deviation of < 2 for better illustration.

```
def random_portfolio(returns):
    '''
    Returns the mean and standard deviation of returns for a random portfolio
    '''

    p = np.asmatrix(np.mean(returns, axis=1))
    w = np.asmatrix(rand_weights(returns.shape[0]))
    C = np.asmatrix(np.cov(returns))

    mu = w * p.T
    sigma = np.sqrt(w * C * w.T)

    # This recursion reduces outliers to keep plots pretty
    if sigma > 2:
        return random_portfolio(returns)
    return mu, sigma
```

In the code you will notice the calculation of the return with:

$$R = p^T w$$

where R is the expected return, p^T is the transpose of the vector for the mean returns for each time series and w is the weight vector of the portfolio. p is a $N \times 1$ column vector, so p^T turns into a $1 \times N$ row vector which can be multiplied with the $N \times 1$ weight (column) vector w to give a scalar result. This is equivalent to the dot product used in the code. Keep in mind that **Python** has a reversed definition of rows and columns and the accurate **NumPy** version of the previous equation would be `$R = w * p.T$`

Next, we calculate the standard deviation with

$$\sigma = \sqrt{w^T C w}$$

where C is the covariance matrix of the returns which is a $N \times N$ matrix. Please note that if we simply calculated the simple standard deviation with the appropriate weighting using `$\text{std}(\text{array}(\text{ret_vec}).T * w)$` we would get a slightly different

[August 2015](#)
[July 2015](#)
[June 2015](#)
[May 2015](#)
[April 2015](#)
[March 2015](#)
[February 2015](#)
[January 2015](#)
[December 2014](#)
[November 2014](#)
[September 2014](#)
[July 2014](#)
[June 2014](#)
[May 2014](#)
[April 2014](#)
[March 2014](#)
[February 2014](#)
[January 2014](#)
[December 2013](#)
[November 2013](#)
[October 2013](#)
[August 2013](#)
[June 2013](#)
[May 2013](#)
[April 2013](#)
[March 2013](#)
[February 2013](#)
[January 2013](#)
[December 2012](#)
[November 2012](#)
[October 2012](#)
[September 2012](#)
[August 2012](#)
[July 2012](#)
[May 2012](#)

Search for:

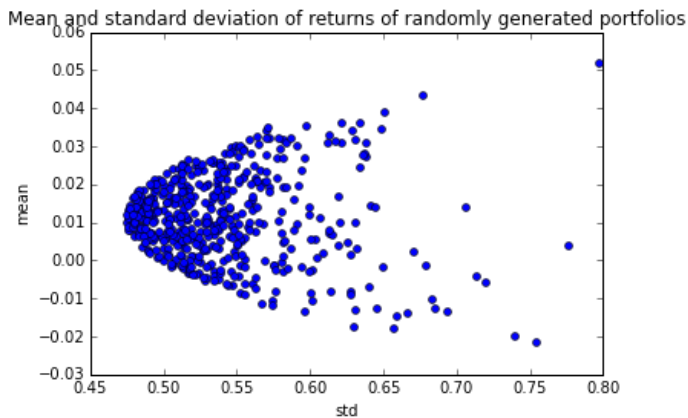
'bullet'. This is because the simple standard deviation calculation would not take covariances into account. In the covariance matrix, the values of the diagonal represent the simple variances of each asset while the off-diagonals are the variances between the assets. By using ordinary `std()` we effectively only regard the diagonal and miss the rest. A small but significant difference.

Lets generate the mean returns and volatility for 500 random portfolios:

```
n_portfolios = 500
means, stds = np.column_stack([
    random_portfolio(return_vec)
    for _ in xrange(n_portfolios)
])
```

Upon plotting those you will observe that they form a characteristic parabolic shape called the 'Markowitz bullet' with the boundaries being called the 'efficient frontier', where we have the lowest variance for a given expected.

```
plt.plot(stds, means, 'o', markersize=5)
plt.xlabel('std')
plt.ylabel('mean')
plt.title('Mean and standard deviation of returns of randomly generated portfolios')
```



Markowitz optimization and the Efficient Frontier

Once we have a good representation of our portfolios as the blue dots show we can calculate the efficient frontier Markowitz-style. This is done by minimising

$$w^T C w$$

for w on the expected portfolio return $R^T w$ whilst keeping the sum of all the weights equal to 1:

$$\sum_i w_i = 1$$

Here we parametrically run through $R^T w = \mu$ and find the minimum variance for different μ 's. This can be done with `scipy.optimize.minimize` but we have to define quite a complex problem with bounds, constraints and a Lagrange multiplier. Conveniently, the `cvxopt` package, a convex solver, does all of that for us. We used one of their [examples](#) with some modifications as shown below. You will notice that there are some conditioning expressions in the code. They are simply needed to set up the problem. For more information please have a look at the `cvxopt` example.

The `mus` vector produces a series of expected return values μ in a non-linear and more appropriate way. We will see later that we don't need to calculate a lot of these as they perfectly fit a parabola, which can safely be extrapolated for higher values.

```
def optimal_portfolio(returns):
    n = len(returns)
    returns = np.asmatrix(returns)

    N = 100
    mus = [10**(5.0 * t/N - 1.0) for t in range(N)]
```

```
# Convert to cvxopt matrices
S = opt.matrix(np.cov(returns))
pbar = opt.matrix(np.mean(returns, axis=1))

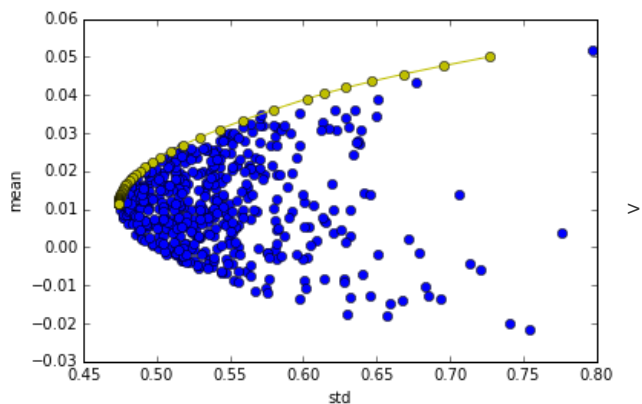
# Create constraint matrices
G = -opt.matrix(np.eye(n)) # negative n x n identity matrix
h = opt.matrix(0.0, (n, 1))
A = opt.matrix(1.0, (1, n))
b = opt.matrix(1.0)

# Calculate efficient frontier weights using quadratic programming
portfolios = [solvers.qp(mu*S, -pbar, G, h, A, b)['x']
               for mu in mus]

## CALCULATE RISKS AND RETURNS FOR FRONTIER
returns = [blas.dot(pbar, x) for x in portfolios]
risks = [np.sqrt(blas.dot(x, S*x)) for x in portfolios]
## CALCULATE THE 2ND DEGREE POLYNOMIAL OF THE FRONTIER CURVE
m1 = np.polyfit(returns, risks, 2)
x1 = np.sqrt(m1[2] / m1[0])
# CALCULATE THE OPTIMAL PORTFOLIO
wt = solvers.qp(opt.matrix(x1 * S), -pbar, G, h, A, b)['x']
return np.asarray(wt), returns, risks

weights, returns, risks = optimal_portfolio(return_vec)

plt.plot(stds, means, 'o')
plt.ylabel('mean')
plt.xlabel('std')
plt.plot(risks, returns, 'y-o')
```



In yellow you can see the optimal portfolios for each of the desired returns (i.e. the `mus`). In addition, we get the one optimal portfolio returned:

```
print weights
```

```
[[ 2.77880107e-09]
 [ 3.20322848e-06]
 [ 1.54301198e-06]
 [ 9.99995251e-01]]
```

Backtesting on real market data ¶

This is all very interesting but not very applied. We next demonstrate how you can create a simple algorithm in `zipline` -- the open-source backtester that powers [Quantopian](#) -- to test this optimization on actual historical stock data.

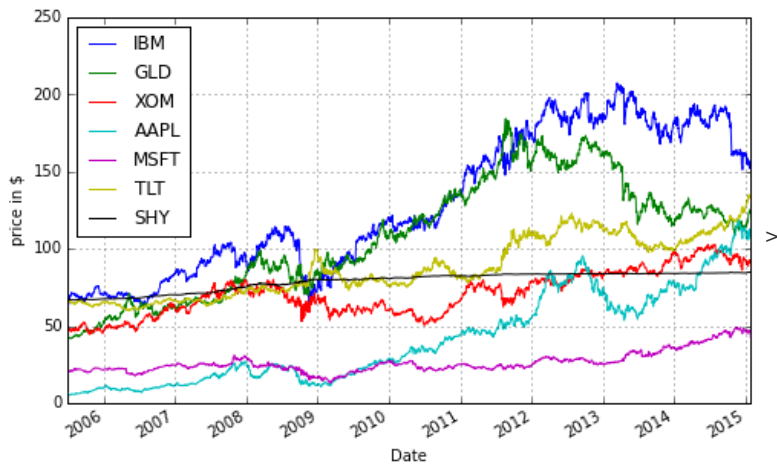
First, let's load in some historical data using [Quantopian's](#) data (if we are running in the [Quantopian Research Platform](#), or the `load_bars_from_yahoo()` function from `zipline`).

```
from zipline.utils.factory import load_bars_from_yahoo
end = pd.Timestamp.utcnow()
start = end - 2500 * pd.tseries.offsets.BDay()

data = load_bars_from_yahoo(stocks=['IBM', 'GLD', 'XOM', 'AAPL',
                                   'MSFT', 'TLT', 'SHY'],
                           start=start, end=end)
```

IBM
GLD
XOM
AAPL
MSFT
TLT
SHY

```
data.loc[:, :, 'price'].plot(figsize=(8,5))
plt.ylabel('price in $')
```



Next, we'll create a `zipline` algorithm by defining two functions -- `initialize()` which is called once before the simulation starts, and `handle_data()` which is called for every trading bar. We then instantiate the algorithm object.

If you are confused about the syntax of `zipline`, check out the [tutorial](#).

```
import zipline
from zipline.api import (add_history,
                        history,
                        set_slippage,
                        slippage,
                        set_commission,
                        commission,
                        order_target_percent)

from zipline import TradingAlgorithm

def initialize(context):
    """
    Called once at the very beginning of a backtest (and live trading).
    Use this method to set up any bookkeeping variables.

    The context object is passed to all the other methods in your algorithm.

    Parameters

    context: An initialized and empty Python dictionary that has been
             augmented so that properties can be accessed using dot
             notation as well as the traditional bracket notation.

    Returns None
    """
    # Register history container to keep a window of the last 100 prices.
    add_history(100, 'ld', 'price')
    # Turn off the slippage model
    set_slippage(slippage.FixedSlippage(spread=0.0))
    # Set the commission model (Interactive Brokers Commission)
    set_commission(commission.PerShare(cost=0.01, min_trade_cost=1.0))
    context.tick = 0

def handle_data(context, data):
    """
    Called when a market event occurs for any of the algorithm's
    securities.
```

Parameters

data: A dictionary keyed by security id containing the current state of the securities in the algo's universe.

context: The same context object from the initialize function. Stores the up to date portfolio as well as any state variables defined.

Returns None

```
'''
# Allow history to accumulate 100 days of prices before trading
# and rebalance every day thereafter.
context.tick += 1
if context.tick < 100:
    return
# Get rolling window of past prices and compute returns
prices = history(100, '1d', 'price').dropna()
returns = prices.pct_change().dropna()
try:
    # Perform Markowitz-style portfolio optimization
    weights, _, _ = optimal_portfolio(returns.T)
    # Rebalance portfolio accordingly
    for stock, weight in zip(prices.columns, weights):
        order_target_percent(stock, weight)
except ValueError as e:
    # Sometimes this error is thrown
    # ValueError: Rank(A) < p or Rank([P; A; G]) < n
    pass

# Instantiate algorithm
algo = TradingAlgorithm(initialize=initialize,
                        handle_data=handle_data)

# Run algorithm
results = algo.run(data)
results.portfolio_value.plot()
```

```
[2015-01-28 14:35:58.352355] INFO: Performance: Simulated 2411 trading days o
ut of 2411.
[2015-01-28 14:35:58.352976] INFO: Performance: first open: 2005-06-29 13:31:
00+00:00
[2015-01-28 14:35:58.353412] INFO: Performance: last close: 2015-01-27 21:00:
00+00:00
```



As you can see, the performance here is pretty good, even through the 2008 financial crisis. This is most likely due to our universe selection and shouldn't always be expected. Increasing the number of stocks in the universe might reduce the volatility as well. Please let us know in the comments section if you had any success with this strategy and how many stocks you used.

Conclusions¶

In this blog, co-written by Quantopian friend [Dr. Thomas Starke](#), we wanted to provide an intuitive and gentle introduction to Markowitz portfolio optimization which still remains relevant today. By using simulation of various random portfolios we have seen that certain portfolios perform better than others. Convex optimization using `cvxopt` allowed us to then numerically determine the portfolios that live on the *efficient frontier*. The zipline backtest serves as an example but also shows compelling performance.


Next steps¶

- Clone this notebook in the [Quantopian Research Platform](#) and run it on your own to see if you can enhance the performance. See [here](#) for a forum thread containing this Notebook on Quantopian.
- You can also download just the notebook for use in your own environment [here](#).
- Read a recent interview with Harry Markowitz: [What Does Harry Markowitz Think?](#)
- In a future blog post we will outline the connections to Kelly optimization which also tells us the amount of leverage to use.
- We are currently in the process of adding `cvxopt` to the Quantopian backtester -- stay tuned!




This entry was posted on Monday, March 9th, 2015 at 12:03 pm and is filed under [Quantopian](#). You can follow any responses to this entry through the [RSS 2.0](#) feed. Both comments and pings are currently closed.


4 responses to “The Efficient Frontier: Markowitz portfolio optimization in Python”

1.  Warren O'Neill says:
March 10, 2015 at 4:31 am


Does anyone have experience with the cvxpy (<https://github.com/cvxgrp/cvxpy>) packet? It seems to be a good alternative to cvxopt.

2.  Thomas Wiecki says:
March 10, 2015 at 7:28 am

Warren: Thanks for your comment. Yeah, I looked into cvxpy and it looks like it'd be much easier to use. It's basically a wrapper for cvxopt and some other optimizers. I think this code could easily be rewritten to use cvxpy instead.

3.  Warren O'Neill says:
March 13, 2015 at 8:08 am

Thanks Thomas! and thanks of the interesting post.

4.  Fred says:
March 27, 2015 at 5:43 am

Hi,
thanks for the nice article, but the link to the cvxopt website does not seem to work....

© 2013-2016 Quantopian

SITE HELP[Frequently Asked Questions](#)[Help & API Docs](#)[Contact Us](#)**MORE INFO**[Our Blog](#)[About Us](#)[Terms of Use](#)[Privacy Policy](#)**FOLLOW US**[Twitter](#)[Facebook](#)[LinkedIn](#)[Blog RSS](#)
